

2.1.10 Live Lab: Create and Deploy an Azure OpenAI LLM

Jeffrey Jordan (ghostinthemaking@gmail.com)

Date: 10:20:31 PM 03/04/2026

Time Spent: 00:51:52

Score: 100%

Passing Score: 80%



Task: Create the Azure OpenAI service

✓ Correct

Task: Deploy the gpt-4o model

✓ Correct

Question 1

✓ Correct

What does max response primarily control?

- ☐ The number of messages in the conversation history
- ☒ The maximum length of the model's next response ✓ Correct
- ☐ The maximum number of requests per minute
- ☐ The sampling temperature

Question 2

✓ Correct

Lowering temperature generally makes responses...

- ☐ Longer
- ☐ More random
- ☒ More deterministic/repeatable ✓ Correct
- ☐ Cheaper

Question 3

✓ Correct

If your AI model's response is being cut off prematurely, which of the following are recommended fixes in this lab's context? (Select all that apply)

- ☒ Increase max response ✓ Correct
- ☒ Reduce the amount of output your prompt requests ✓ Correct
- ☐ Increase the sampling temperature
- ☐ Switch to a smaller model

Question 4

✓ Correct

What is a potential downside of setting max response much higher than needed?

- ☒ Increased latency and cost per request ✓ Correct
- ☐ Lower quality responses
- ☐ The model becomes unable to process shorter prompts
- ☐ Requests are automatically rejected by Azure AI

