

4.2.8: Live Lab: Test Prompt Injection Attacks

SecAI+ (Exam CY0-001) - PROD

Duration: 1 hour, 19 minutes



Congratulations, you passed!

Group 1

Score: 4/4, 100%

☒ Role-play attacks are most likely to succeed when: Score: 1

- ☐ The model has strong instruction precedence and detects roleplay attempts
- ☐ The system prompt is restrictive and enforces strict boundaries
- ☒ The model readily accepts persona instructions or the system prompt is permissive
- ☐ Output filtering is enabled on all responses

Correct! Role-play attacks work when the model readily accepts persona instructions or roleplay, especially when the system prompt is permissive or the model heavily weights recent user role instructions.

☒ What is the main goal of a justification extraction attack? Score: 1

- ☐ To force the model to encode outputs in base64
- ☒ To ask the model to explain why it refused a request, potentially revealing internal rules
- ☐ To make the model adopt an admin persona
- ☐ To inject malicious instructions into indexed documents

Correct! Justification extraction asks the model to explain or justify a prior refusal, turning a safety explanation into a vector for leaking hidden prompts or rules that attackers can then exploit.

☒ In output coercion attacks, why do attackers request specific formats like base64 encoding? Score: 1

- ☐ To make the response more readable for end users
- ☐ To improve model performance and accuracy
- ☒ To hide literal keywords and slip past output filters while making exfiltration easier
- ☐ To compress large responses into smaller outputs

Correct! Encoding hides literal keywords and can slip past output filters that search for known token formats; encoded strings are also easier to exfiltrate and later decode.

☒ What distinguishes indirect prompt injection from direct prompt injection?

Score: 1

- ☐ It only works against newer language models
- ☒ Malicious instructions are hidden in data the system fetches or indexes
- ☐ It requires physical access to the server
- ☐ It can only be performed by system administrators

Correct! Indirect prompt injection occurs when attacker-controlled instructions are hidden inside data the system fetches or indexes (documents, web pages, uploaded files), reaching the model via the retriever rather than direct chat input.

☒ **Check Prompts 1**

Score: 3

Prompts validated!

☒ **Check Prompts 2**

Score: 3

Prompts validated!

☒ Successful prompt(s):

Not answered