6.2.7 Lesson Review

Date: 4/7/2026, 3:54:37 PM

Time Spent: 08:08

Score: 100%

Passing Score: 80%

Question 1 Correct

How can the output of an AI system most directly lead to reputational harm for an organization?

- ○ The AI provides technically correct information that some users misinterpret or find confusing.
- ○ The AI output is reviewed by staff, so any reputational impact is usually identified and addressed.
- ○ The AI produces content that conflicts with the organization's stated values or policies.
- ◉ The AI generates offensive, inaccurate, or inappropriate content that reaches customers. ✓ Correct

Explanation

The AI generates offensive, inaccurate, or inappropriate content that reaches customers is the correct answer. When AI outputs are misleading or offensive, it can damage customer trust and the organization's reputation, regardless of intent.

The AI produces content that conflicts with the organization's stated values or policies is incorrect. This can cause issues, but the harm is more indirect and often depends on interpretation, compared to clearly offensive or inaccurate content that directly reaches customers.

The AI provides technically correct information that some users misinterpret or find confusing is incorrect. Confusing but accurate content may frustrate users, yet it is less likely to cause immediate, widespread reputational harm than clearly offensive or wrong output.

The AI output is reviewed by staff, so any reputational impact is usually identified and addressed is incorrect. This describes a control to reduce risk, not a primary way reputational harm occurs. The main risk still comes from harmful content that reaches customers.

Related Content

6.2.4 Common AI Risks

resources\questions\q_common_ai_risks_03.question.xml

Question 2 ✓ Correct

A healthcare company is preparing to launch an AI model to help therapists suggest customized treatment plans.

To support responsible AI governance, what is the BEST way to document the system before deployment?

- ○ Rely on verbal explanations during training sessions to communicate model information, updating the content only if staff raise questions or concerns.
- ○ List only the names of the development team and the project timeline, since detailed technical information may confuse non-technical stakeholders.
- ○ Prepare a general user manual describing how to access the AI system, but omit details on data sources, intended users, or system training.
- ◉ Create a concise model card outlining the AI system's purpose, data sources, training period, limitations, intended users, and prohibited uses. ✓ Correct

Explanation

Creating a concise model card outlining the AI system's purpose, data sources, training period, limitations, intended users, and prohibited uses is the correct answer. It follows best practice by giving structured information on purpose, boundaries, data, and limits, improving transparency and auditability.

Listing only the development team and project timeline is incorrect. It omits key details such as data, usage restrictions, and intended function.

Preparing a general user manual on how to access the system but omitting data, intended users, or training details is incorrect. It lacks necessary context for safe use.

Relying on verbal explanations during training, updated only when questions arise, is incorrect. Verbal-only info is inconsistent, untracked, and noncompliant.

Related Content

resources\questions\q_putting_principles_into_practice_02.question.xml

Question 3 Correct

Which of the following BEST describes a key intellectual property (IP) risk associated with using AI models trained on external data sources?

- Training data collected from publicly accessible websites is generally considered low risk for copyright issues and can typically be reused in AI models with minimal legal concern.
- AI models trained on user data can be reused across multiple unrelated projects without performing any kind of data source audit to certify legal compliance, provided the data is de-identified.
- The AI system may inadvertently generate outputs that replicate or closely mimic copyrighted or proprietary materials. ✓ Correct
- The AI model may operate with reduced accuracy if deployed in changing business environments or with new data types.

Explanation

The AI system may inadvertently generate outputs that replicate or closely mimic copyrighted or proprietary materials is the correct answer. AI can unintentionally reproduce protected content, creating legal and financial risks.

The AI model may operate with reduced accuracy if deployed in changing business environments or with new data types is incorrect. This describes a performance issue (model drift), not an IP or provenance risk.

Training data collected from publicly accessible websites is generally considered low risk for copyright issues is incorrect. Public accessibility does not remove copyright or licensing restrictions. Many online sources are still copyrighted or governed by terms of use, and using them for AI training can require permissions or licenses

Reusing trained AI model across unrelated projects without data source audits if data is de-identified is incorrect. Even de-identified user data requires compliance checks and licenses; reckless reuse can violate intellectual property rights.

Related Content

6.2.4 Common AI Risks

resources\questions\q_common_ai_risks_05.question.xml

Question 4 Correct

A cybersecurity company uses an AI model to detect phishing emails. An attacker manages to inject crafted emails into the training dataset so the model incorrectly classifies certain phishing attempts as safe.

The security team is concerned about data poisoning and seeks ways to reduce this risk in future AI training cycles.

Which action BEST applies a risk mitigation step for data poisoning as described in the scenario?

- Assign one team member to subjectively hand-pick training examples and exclude suspicious emails.
- Implement integrity checks and isolate the training pipeline to ensure only verified and safe data is used. ✓ Correct
- Retrain the model more frequently, regardless of changes in the training data.
- Ignore suspicious activity in the training data unless there's a known, large-scale breach.

Explanation

Implementing integrity checks and isolating the training pipeline to ensure only verified and safe data is used is the correct answer. It directly addresses data poisoning by verifying input trustworthiness and preventing tampering, reducing the chance that malicious samples influence the model.

Retraining the model more frequently, regardless of changes in the training data, is incorrect. If the pipeline is compromised, retraining can reinforce poisoned patterns. Protection requires securing and validating data, not just retraining.

Assigning one team member to hand-pick training examples and exclude suspicious emails is incorrect. This approach is subjective, error-prone, and easily bypassed by subtle manipulations. Robust, automated integrity checks with controlled isolation provide stronger protection.

Ignoring suspicious activity unless there's a known, large-scale breach is incorrect. Early signs must be investigated.

Related Content

resources\questions\q_identifying_risks_unique_to_ai_05.question.xml

Question 5 Correct

A financial institution is developing an AI-powered loan approval system. During testing, the team observes that applicants from one demographic group are consistently approved at a significantly higher rate than those from another, despite having similar credit profiles.

Based on the principles of responsible AI with an emphasis on fairness, which action should the team take to address this issue?

- Keep the model as is since it meets standard technical accuracy metrics and has already passed initial security testing.
- Rely on customer service staff to manually override unfair decisions when detected, rather than adjusting the model or data.
- Conduct a thorough review of the training data, implement fairness metrics to measure approval rates across groups, and adjust the model or data to correct any unjustified disparities. ✓ Correct
- Increase the size of the training dataset by collecting more data from all applicants, but take no other actions regarding group disparities.

Explanation

Conducting a thorough review of the training data, implementing fairness metrics to measure approval rates across groups, and adjusting the model or data to correct unjustified disparities is the correct answer. It directly addresses bias linked to protected characteristics through analysis and corrective action.

Keeping the model as is because it meets accuracy and security metrics is incorrect. Responsible AI requires fairness, not just performance and security.

Increasing the dataset size without addressing group disparities is incorrect. More data alone can preserve or amplify existing bias.

Relying on customer service staff to manually override unfair decisions is incorrect. Human oversight helps, but it does not fix systemic bias in the model and leads to inconsistent, inefficient outcomes.

Related Content

6.2.1 Principles of Responsible AI

resources\questions\q_principles_of_responsible_ai_09.question.xml

Question 6 ✓ Correct

Which of the following BEST reflects an important consideration for sustainability in the development and deployment of AI systems?

- ○ Scaling AI workloads across multiple data centers to balance performance and maintain consistent service availability.
- ○ Increasing hardware utilization through higher model throughput to maximize the return on existing infrastructure investments.
- ○ Optimizing AI infrastructure to minimize operational costs and allocate more budget to future system enhancements.
- ◉ Monitoring energy consumption and using efficient models to reduce the environmental impact of AI operations. ✓ Correct

Explanation

Monitoring energy consumption and using efficient models to reduce the environmental impact of AI operations is the correct answer. Sustainability in AI focuses on the energy and hardware used in training and deployment; efficient models reduce carbon footprint and support environmental responsibility.

Optimizing AI infrastructure to minimize operational costs is incorrect. This option sounds similar to sustainability because it mentions optimization and efficiency, but its primary goal is financial—reducing costs and funding enhancements—rather than explicitly reducing environmental impact.

Scaling AI workloads across multiple data centers is incorrect. This focuses on performance and reliability through scaling and multiple data centers. While operationally important, it does not address tracking or reducing energy use or emissions, so it does not directly support environmental sustainability.

Increasing hardware utilization through higher model throughput is incorrect. This highlights higher utilization and throughput, which can seem efficient, but the emphasis is on maximizing return on investment. It may even increase total energy consumption and does not mention reducing environmental impact, so it is not a strong sustainability measure.

Related Content

resources\questions\q_principles_of_responsible_ai_07.question.xml

Question 7 Correct

How can organizations minimize the risk of accidental data leakage and privacy exposure in AI systems?

- By anonymizing data during storage but keeping full, identifiable records in model training logs.
- By relying on encryption at rest and in transit without restricting who can access the data.
- By using access controls and encryption while relying on manual review instead of automated redaction.
- By using access controls, encryption, and redaction during data processing. ✓ Correct

Explanation

By using access controls, encryption, and redaction during data processing is the correct answer. Effective data governance limits access to sensitive data, encrypts information, and redacts personal data from records or outputs to prevent leaks.

By using access controls and encryption while relying on manual review instead of automated redaction is incorrect. While access controls and encryption are important, relying on manual review for redaction introduces human error and inconsistency, especially at scale. Automated redaction processes are typically needed to systematically remove sensitive information from large volumes of data used by AI systems.

By relying on encryption at rest and in transit without restricting who can access the data is incorrect. Encryption protects data from external interception, but it does not address internal misuse or overexposure. If access is not restricted, anyone within the organization with basic credentials might still view sensitive information, increasing the risk of accidental leakage.

By anonymizing data during storage but keeping full, identifiable records in model training logs is incorrect. Anonymizing stored data helps, but keeping full identifiable information in training logs undermines that protection. Logs can be accessed, misconfigured, or breached, and if they contain identifiable data, they become a significant source of privacy exposure.

Related Content

6.2.1 Principles of Responsible AI

resources\questions\q_principles_of_responsible_ai_05.question.xml

Question 8 ✓ Correct

A financial institution deploys a deep neural network to assess loan applications. The model is highly accurate, but internal auditors raise concerns that its complex structure makes it difficult to interpret why some applicants are denied, and there is a risk of undetected discrimination.

The bank wants to ensure ethical and regulatory compliance while maintaining strong model performance.

Which is the MOST appropriate step to help address the unique risks associated with using a complex AI model in this situation?

- (selected) Apply feature importance analysis to help auditors understand which factors most influence the model's decisions. ✓ Correct
- Ignore the explainability concern since high accuracy suggests the model is working as intended.
- Increase the complexity of the model by adding more layers to improve its accuracy.
- Remove all data fields except the applicant's name to make the model simpler and avoid discrimination.

Explanation

Applying feature importance analysis to help auditors understand which factors most influence the model's decisions is the correct answer. Feature importance highlights key variables affecting predictions, supporting transparency, fairness, and compliance.

Increasing the model's complexity by adding more layers to improve accuracy is incorrect. More layers can reduce interpretability and don't address auditors' concerns about transparency or potential bias.

Removing all data fields except the applicant's name to simplify the model is incorrect. This makes the model ineffective, as names alone provide no insight into credit risk and eliminate model utility.

Ignoring explainability concerns because of high accuracy is incorrect. Accuracy does not guarantee fairness; explainability is essential for ethical, transparent, and regulatory-compliant decision-making.

Related Content

resources\questions\q_identifying_risks_unique_to_ai_04.question.xml

Question 9 Correct

Which approach best helps organizations manage the risk of accuracy decline in deployed AI models?

- Retrain the AI model on updated data according to a fixed schedule.
- Periodically evaluate the AI model against representative validation datasets.
- Continuously monitor live model outputs and compare them to original baselines. ✓ Correct
- Conduct post-deployment testing using a sample of real-world prediction cases.

Explanation

Continuously monitoring live model outputs and comparing them to original baselines is the correct answer. Ongoing monitoring detects shifts in accuracy, allowing organizations to act before major problems arise.

Periodically evaluating the AI model against representative validation datasets is incorrect. This is a good practice and can detect drift, but it happens at intervals and may miss rapid or subtle changes between evaluations. Continuous live monitoring gives earlier, more consistent detection of accuracy decline.

Retraining the AI model on updated data according to a fixed schedule is incorrect. Scheduled retraining can help keep models current, but without performance-based triggers it can be too late or unnecessary. Continuous monitoring lets organizations decide when to retrain based on actual evidence of declining accuracy.

Conducting post-deployment testing using a sample of real-world prediction cases is incorrect. Testing with real cases after deployment is useful, but if it is done only occasionally and on limited samples, important changes can be missed. Continuous monitoring of broad live outputs provides more complete and timely visibility into accuracy issues.

Related Content

6.2.4 Common AI Risks

resources\questions\q_common_ai_risks_04.question.xml

Question 10 Correct

A large retail company introduces an AI-driven recommendation engine that uses customer purchase history, browsing activity, and personal information to suggest products.

After a customer complains that their private information appeared in a product suggestion email, the company's data privacy team reviews its AI practices.

According to the principles of privacy and data governance for responsible AI, what is the BEST step the company should take?

- Temporarily pause the recommendation engine and assign the issue to the marketing team, asking them to handle customer communications and privacy safeguards as needed, before resuming normal operations.
- Delete all customer data from the system and disable the recommendation engine to completely eliminate the risk of privacy exposure.
- Continue collecting and using detailed personal information in the AI system, but inform customers that their data will be used for personalized marketing messages.
- Implement clear data governance controls, use data minimization and de-identification techniques, review data retention policies, and ensure personal data is not included in AI-generated outputs. ✓ Correct

Explanation

Implementing clear data governance controls, using data minimization and de-identification, reviewing retention policies, and ensuring personal data is not in AI outputs is the correct answer. It protects customer information and prevents further breaches.

Continuing to use detailed personal data while informing customers it's for personalized marketing is incorrect. Notification alone, without limiting or protecting data, violates responsible governance.

Deleting all customer data and disabling the recommendation engine is incorrect. It removes risk but is excessive; responsible AI aims for minimal, secure data use while preserving business value.

Temporarily pausing the engine and assigning the issue to marketing is incorrect. Privacy safeguards should be handled through formal governance, not ad hoc by unrelated teams.

Related Content

6.2.1 Principles of Responsible AI

resources\questions\q_principles_of_responsible_ai_12.question.xml