

6.4 Module Quiz

Date: 4/13/2026, 4:11:13 AM

Time Spent: 12:30

Score: 100%

Passing Score: 80%



Which of the following BEST describes a key intellectual property (IP) risk associated with using AI models trained on external data sources?

- ☒ The AI system may inadvertently generate outputs that replicate or closely mimic copyrighted or proprietary materials. ✓ Correct
- ☐ Training data collected from publicly accessible websites is generally considered low risk for copyright issues and can typically be reused in AI models with minimal legal concern.
- ☐ AI models trained on user data can be reused across multiple unrelated projects without performing any kind of data source audit to certify legal compliance, provided the data is de-identified.
- ☐ The AI model may operate with reduced accuracy if deployed in changing business environments or with new data types.

Explanation

The AI system may inadvertently generate outputs that replicate or closely mimic copyrighted or proprietary materials is the correct answer. AI can unintentionally reproduce protected content, creating legal and financial risks.

The AI model may operate with reduced accuracy if deployed in changing business environments or with new data types is incorrect. This describes a performance issue (model drift), not an IP or provenance risk.

Training data collected from publicly accessible websites is generally considered low risk for copyright issues is incorrect. Public accessibility does not remove copyright or licensing restrictions. Many online sources are still copyrighted or governed by terms of use, and using them for AI training can require permissions or licenses

Reusing trained AI model across unrelated projects without data source audits if data is de-identified is incorrect. Even de-identified user data requires compliance checks and licenses; reckless reuse can violate intellectual property rights.

Related Content



6.2.4 Common AI Risks

resources\questions\q_common_ai_risks_05.question.xml

Why is explainability particularly important for high-risk AI applications, such as medical diagnosis or credit assessment?

- ☐ Because it speeds up the AI's decision-making process.
- ☐ Because explainable models always produce higher accuracy than complex, less interpretable models.
- ☒ Because stakeholders need to understand the rationale behind decisions that carry significant consequences. ✓ Correct
- ☐ Because it eliminates the need for any human review or oversight.

Explanation

Because stakeholders need to understand the rationale behind decisions that carry significant consequences is the correct answer. In high-stakes domains, explainability lets users and affected individuals see why decisions were made, improving trust, enabling error correction, and supporting ethical accountability.

Because it speeds up the AI's decision-making process is incorrect. Explainability aids debugging and trust but does not directly make computation faster.

Because it eliminates the need for any human review or oversight is incorrect. Explainability supports, but never replaces, human oversight in high-risk uses.

Because explainable models always produce higher accuracy than complex, less interpretable models is incorrect. Explainable models may trade some accuracy for transparency; the priority here is understandability for critical decisions.

Related Content



6.2.1 Principles of Responsible AI

resources\questions\q_principles_of_responsible_ai_04.question.xml

Question 3

✓ Correct

What does the principle of fairness require organizations to do when developing and deploying AI systems?

- ☐ Ignore group-level performance differences if the system works well for the majority.
- ☒ Establish clear rules for acceptable disparities and take corrective actions if results are biased. ✓ Correct
- ☐ Allow AI systems to learn independently without human intervention.
- ☐ Automatically trust AI outputs as objective and unbiased.

Explanation

Establishing clear rules for acceptable disparities and taking corrective actions if results are biased is the correct answer. Responsible AI requires defined fairness criteria and thresholds; when results exceed these, inequities must be addressed.

Allowing AI systems to learn independently without human intervention is incorrect. Fairness requires proactive oversight to detect and correct imbalances, not hands-off learning.

Automatically trusting AI outputs as objective and unbiased is incorrect. AI can reflect or amplify existing biases if not regularly verified and managed.

Ignoring group-level performance differences if the system works well for the majority is incorrect. Overlooking disparities for minority or protected groups undermines fairness; equitable performance across groups is essential.

Related Content



6.2.1 Principles of Responsible AI

resources\questions\q_principles_of_responsible_ai_02.question.xml

What is a key way in which accidental data leakage can occur in AI systems?

- ☐ AI models only use data from public sources, so leakage isn't possible.
- ☐ AI systems always ignore any personal information in their training data.
- ☒ Sensitive information from training data unintentionally appears in the AI's outputs. ✓ Correct
- ☐ Data leakage can only happen if someone hacks into the system intentionally.

Explanation

Sensitive information from training data unintentionally appears in the AI's outputs is the correct answer. AI systems can expose private data if training data includes sensitive details that are later reproduced in generated content.

AI systems always ignore any personal information in their training data is incorrect. AI does not inherently ignore personal information unless explicitly programmed to filter or redact it.

Data leakage can only happen if someone hacks into the system intentionally is incorrect. Accidental leakage often occurs without a deliberate attack—model outputs can expose unintended data due to weak controls or lack of filtering, even in secure environments.

AI models only use data from public sources, so leakage isn't possible is incorrect. AI systems may use private or internal data, and accidental data leakage is possible even if not using public sources, especially if controls and checks are not in place.

Related Content



6.2.4 Common AI Risks

resources\questions\q_common_ai_risks_02.question.xml

An insurance organization recently adopted AI to automate claims processing. After a compliance audit, the company discovered gaps: several AI models had been deployed without the required approval evidence, records of data sources and model lineage were incomplete, and retention policies set in the company's AI procedures were not consistently enforced.

Leadership requests the AI team to identify the root causes and ensure these governance failures are prevented in the future by integrating policy requirements into day-to-day operations using automated workflows and monitoring systems.

Which team member should take the lead in analyzing these governance challenges and integrating controls to automate approval workflows, manage deployment risks, ensure evidence collection, and enforce data retention and audit requirements?

- ☐ AI Auditor
- ☐ Platform Engineer
- ☐ Data Scientist
- ☒ AI Governance Engineer ✓ Correct

Explanation

AI Governance Engineer is the correct answer. The AI Governance Engineer specializes in operationalizing policies by embedding approval workflows, collecting evidence automatically, and enforcing retention and audit requirements. This role ensures policies are applied consistently throughout the AI lifecycle.

Data Scientist is incorrect. While Data Scientists may document their work, their primary focus is on model development and analysis, not the automation and integration of governance controls and evidence collection as required for comprehensive policy compliance.

Platform Engineer is incorrect. Although Platform Engineers build infrastructure and may deploy monitoring tools, they do not integrate governance controls and automate approval workflows according to detailed policy requirements.

AI Auditor is incorrect. The AI Auditor is responsible for creating evidence collection systems and configuring policy-driven approval workflows directly within operational AI platforms to prevent governance issues. AI Auditors evaluate compliance and recommend improvements but do not design or embed automated governance processes or controls.

Related Content



6.1.2 Important Roles in AI

resources\questions\q_important_roles_in_ai_08.question.xml

Question 6

✓ Correct

Match each AI-related role on the left with its correct description on the right.

Builds and maintains data pipelines to acquire, clean, label, and deliver quality-controlled data for AI initiatives.

✓ Data Engineer

Explores and analyzes data to define problems, select features, and develop baseline AI models that show value.

✓ Data Scientist

Transforms prototype models into production-grade systems by optimizing code, packaging dependencies, and constructing services for reliable, real-world operation.

✓ Machine Learning Engineer

Designs safeguards to protect AI systems from threats such as data leakage, model poisoning, and prompt injection, ensuring that data, models, and prompts are secure.

✓ AI Security Architect

Provides independent assurance by reviewing artifacts and logs, verifying policy adherence, and confirming that monitoring and incident handling are operating as intended.

✓ AI Auditor

Explanation

The following are the correct description matches to the listed AI-related roles:

1. **Data Engineer.** The data engineer's primary responsibility is to build and maintain the complex pipelines that gather, prepare, and ensure the quality of data used in AI systems. This foundation is crucial for trustworthy and reliable AI modeling.
2. **Data Scientist.** The data scientist explores the data provided, determines the relevant features, clarifies problems, and crafts initial or baseline AI models—these early models demonstrate the possible value of more robust solutions.
3. **Machine Learning Engineer.** Machine learning engineers take the experimental models created by data scientists and turn them into robust, scalable products. This involves code optimization, dependency management, and designing production-ready inference services.
4. **AI Security Architect.** This role is centered on risk and protection. The AI security architect devises and implements measures to secure data, models, and processes from a range of possible attacks and vulnerabilities using technical and procedural safeguards.
5. **AI Auditor.** The AI auditor ensures the whole governance process works as intended—reviewing artifacts, logs, and documentation to verify that governance policies are consistently followed, and that incident response and monitoring are functioning as required.

Related Content

resources\questions\q_important_roles_in_ai_01.question.xml

A financial services company is deploying an AI model to automate loan approvals. Management is concerned about potential threats, such as malicious attempts to manipulate model inputs (prompt injection), leakage of personally identifiable information, unauthorized access to sensitive model data, and the possibility of biased outputs affecting regulatory compliance.

The company asks the AI team to examine these risks and establish targeted measures for prevention and detection.

Which team member is BEST positioned to analyze these threats and design integrated security controls such as secrets management, access restrictions, prompt and content filtering, and adversarial testing to ensure that the AI model operates securely and in compliance with financial regulations?

☒ AI Security Architect ✓ Correct

☐ Platform Engineer

☐ Data Scientist

☐ AI Risk Analyst

Explanation

AI Security Architect is the correct answer. The AI Security Architect is specifically tasked with evaluating the full range of threats to AI systems and designing technical safeguards, such as secrets management and adversarial testing, to address these risks. This role is the one responsible for integrating comprehensive protections at every layer of the AI system.

Data Scientist is incorrect. While Data Scientists are critical for model building and testing, their primary focus is on model accuracy, relevance, and performance—not the systematic design and management of security controls.

Platform Engineer is incorrect. Platform Engineers enable the infrastructure for AI development and deployment, but the specific analysis of security risks and design of protective measures against threats such as prompt injection and data leakage is outside their primary scope.

AI Risk Analyst is incorrect. The AI Risk Analyst evaluates risks and recommends mitigations but does not directly create technical security controls.

Related Content



6.1.2 Important Roles in AI

resources\questions\q_important_roles_in_ai_07.question.xml

A U.S.-based startup offers an AI-powered hiring platform to companies in Germany and France. The system screens résumés and automatically rejects candidates who don't meet certain criteria.

The startup has no office in Europe and assumed EU laws would not apply. A German client warns them that their system is considered "high-risk" under the EU AI Act.

Which action BEST aligns with the startup's obligations under the EU AI Act?

- ☒ Classify the hiring system as high-risk, implement required controls (documentation, data governance, testing, human oversight), and complete the conformity assessment for access to the EU market. ✓ Correct
- ☐ Keep operating as usual because the startup is not physically located in the EU, and the EU AI Act only applies to companies headquartered in EU Member States.
- ☐ Stop offering the platform to EU customers and keep all current EU client data, since the Act only applies to future deployments of AI systems.
- ☐ Implement a full AI Management System based on ISO/IEC 42001 and publish an internal AI code of ethics for employees.

Explanation

Classifying the hiring system as high-risk, implementing required controls (documentation, data governance, testing, human oversight), and completing the conformity assessment for access to the EU market is the correct answer. The EU AI Act would apply to this scenario and categorize this as high risk. AI systems used in employment decisions, such as automated candidate screening and rejection, are categorized as high-risk because they significantly affect individuals' rights and opportunities. For high-risk AI, providers must implement extensive controls and put in place post-market monitoring. They must also undergo a conformity assessment before placing the system on the EU market.

Implementing a full AI Management System based on ISO/IEC 42001 and publishing an internal AI code of ethics for employees is incorrect. The question asks what best aligns with obligations specifically under the EU AI Act. The Act requires risk-based categorization, technical and organizational measures, documentation, and conformity assessments for high-risk systems. Merely adopting ISO/IEC 42001 does not ensure compliance with the Act's requirements. This step alone does not adequately address the startup's regulatory responsibilities for placing a high-risk AI system on the EU market and would leave significant legal gaps under the EU AI Act.

Stopping offering the platform to EU customers and keep all current EU client data, since the Act only applies to future deployments of AI systems, is incorrect. Ceasing future sales might reduce ongoing exposure but the startup still has existing EU customers actively using the hiring platform, so compliance duties under the EU AI Act remain for those deployments (such as documentation and post-market monitoring). Additionally, retaining all EU client data without reviewing legal bases and data governance may introduce further regulatory risks, including under other EU laws.

Keeping operating as usual because the startup is not physically located in the EU is incorrect. The EU AI Act applies extraterritoriality to providers that place AI systems on the EU market or whose systems are used within the EU, regardless of where the company is headquartered. The startup offers its AI hiring platform to customers in Germany and France, meaning the system is effectively placed on the EU market and used in Member States. This choice ignores the Act's fundamental structure and would expose the company to potentially severe penalties for noncompliance.

Related Content



6.3.2 Important AI Compliance Frameworks

resources\questions\q_important_ai_compliance_frameworks_05.question.xml

A regional hospital is drafting discharge instructions that include details about a patient's diagnosis, medications, and follow-up appointments.

The hospital's AI policy states that any system processing identifiable health information must use private models in an isolated environment with customer-managed encryption keys and must not allow the provider to train on that data.

A project manager is deciding how to pilot an AI assistant to help clinicians generate discharge instructions.

Which option BEST applies the organization's AI policy on private vs. public models?

- ☒ Deploy a private model in the hospital's virtual private cloud with encryption keys managed by the hospital and prompts excluded from provider training. ✓ Correct
- ☐ Use a popular public AI model with a free web interface and add a banner reminding staff not to paste "too much" patient data.
- ☐ Use public models for all health data but sign a nondisclosure agreement with the AI vendor.
- ☐ Use a public AI model but remove patient names and addresses before pasting discharge notes into the chat.

Explanation

Deploying a private model in the hospital's virtual private cloud with encryption keys managed by the hospital and prompts excluded from provider training is the correct answer. This option correctly applies the policy's requirements for high-risk data. A private model in a VPC provides logical isolation, better control over data residency, and stronger guarantees that PHI (protected health information) doesn't leave the organization's trust boundary. Customer-managed keys and disabled provider training address data-use concerns. For sensitive health data, private models with explicit isolation and contractual controls are preferred over public models, aligning with the policy's guidance.

Using a popular public AI model with a free web interface and adding a banner reminding staff not to paste "too much" patient data is incorrect. This still relies on a public model whose data handling, logging, and training practices are not controlled by the hospital. A banner is not a technical safeguard and is unlikely to consistently prevent PHI from being pasted. The policy explicitly requires private models in an isolated environment for identifiable health information. This choice ignores the need for data residency control, contractual guarantees, and strict configuration and amounts to shadow AI, not compliant use.

Using a public AI model but removing patient names and addresses before pasting discharge notes into the chat is incorrect. Redacting obvious identifiers is insufficient for this scenario. Discharge instructions often contain combinations of dates, diagnoses, procedures, and locations that can still re-identify patients. Even partially anonymized PHI sent to a public model may be logged or used for training outside the hospital's control.

Using public models for all health data but signing a nondisclosure agreement with the AI vendor is incorrect. An NDA does not guarantee technical isolation, data residency, encryption control, or training restrictions. Public models often involve multi-tenant infrastructure and provider-controlled logging, which conflict with the hospital's requirement for private, isolated environments. The policy requires structural controls, not just legal promises, when handling identifiable health information.

Related Content



6.3.3 Organizational AI Policies

resources\questions\q_organizational_ai_policies_05.question.xml

A tech startup is developing a large AI model for real-time language translation, but the engineering team is concerned about the high electricity consumption and environmental costs associated with training and running the model.

To align with the principles of sustainability and social impact for responsible AI, what should the team do?

- ☒ Evaluate whether a smaller, fine-tuned model can deliver similar translation quality while using fewer resources, and monitor the environmental impact of model operations. ✓ Correct
- ☐ Increase AI training frequency to ensure continuous improvements, justifying energy use by potential profits, and wait until environmental regulations require changes.
- ☐ Prioritize translation accuracy as the primary goal and address sustainability in a later development phase.
- ☐ Outsource all AI training to cloud providers, relying on their sustainability statements without conducting an internal assessment of resource usage and impact.

Explanation

Evaluating whether a smaller, fine-tuned model can deliver similar translation quality with fewer resources, and monitoring environmental impact, is the correct answer. It supports sustainability by reducing resource use while maintaining effectiveness.

Prioritize translation accuracy as the primary goal and address sustainability in a later development phase is incorrect. Responsible AI requires integrating environmental and social impact considerations throughout the design and development process, not deferring them until after user adoption.

Increasing training frequency, justifying energy use by profit and waiting for regulation, is incorrect. It prioritizes profit over proactive impact reduction and can greatly increase emissions.

Outsourcing all training to cloud providers and relying only on their sustainability claims is incorrect. Organizations must assess and monitor their own resource use and impacts, not delegate responsibility entirely to vendors.

Related Content

resources\questions\q_principles_of_responsible_ai_14.question.xml

Why does the EU AI Act represent a shift from best practice guidance to mandatory compliance for organizations operating in or selling to the EU market?

- ☐ Because it replaces all existing national and international standards, eliminating the need for organizations to follow frameworks such as ISO or the NIST AI RMF.
- ☒ Because it imposes legally binding, risk-based requirements and steep financial penalties for noncompliance, making adherence to its standards obligatory rather than optional. ✓ Correct
- ☐ Because it applies only to AI research laboratories and exempts commercial providers, allowing them to choose whether to follow the Act based on their internal policies.
- ☐ Because it encourages organizations to invest in trustworthy AI and develop enabling ecosystems, without specifying enforceable obligations or defined risk tiers.

Explanation

Because it imposes legally binding, risk-based requirements and steep financial penalties for noncompliance, making adherence to its standards obligatory rather than optional is the correct answer. The EU AI Act is binding legislation from 2024, with most obligations applying from 2026. It classifies AI by risk tier and sets concrete obligations for high- and limited-risk systems, including documentation, governance, testing, and transparency. Member States must enforce fines, including very steep penalties. This moves organizations from treating responsible AI as discretionary "best practice" to treating compliance as a legal requirement for accessing or serving the EU market.

Because it replaces all existing national and international standards, eliminating the need for organizations to follow frameworks such as ISO or the NIST AI RMF is incorrect. The Act does not abolish or replace other standards and frameworks. In practice, organizations often use ISO standards and frameworks to operationalize compliance and demonstrate maturity. The EU AI Act sets regulatory requirements, while standards and frameworks provide methods for meeting them.

Because it merely encourages organizations to invest in trustworthy AI and develop enabling ecosystems, without specifying enforceable obligations or defined risk tiers is incorrect. This answer confuses the EU AI Act with the OECD AI Principles. The OECD Principles encourage investments in trustworthy AI and enabling ecosystems but are not legally binding.

Because it applies only to AI research laboratories and exempts commercial providers, allowing them to choose whether to follow the Act based on their internal policies is incorrect. The EU AI Act is not limited to research labs nor does it exempt commercial providers. It explicitly governs providers, deployers, importers, and distributors of AI systems in the EU.

Related Content



6.3.2 Important AI Compliance Frameworks

resources\questions\q_important_ai_compliance_frameworks_01.question.xml

How do ISO/IEC 23894:2023, ISO/IEC 22989:2022, and ISO/IEC 5338:2023 collectively support organizations using AI?

- ☒ They provide AI-specific risk management guidance, shared terminology, and lifecycle processes, helping integrate AI risks into enterprise programs and reduce confusion across teams. ✓ Correct
- ☐ They are primarily marketing-focused documents that standardize how AI vendors should describe product features and pricing tiers to enterprise buyers.
- ☐ They function solely as consumer protection laws, focusing on deceptive AI advertising and providing no guidance on technical lifecycle or risk processes.
- ☐ They collectively establish a single binding EU regulation that bans unacceptable-risk AI applications and imposes steep penalties for noncompliance across all Member States.

Explanation

Providing AI-specific risk management guidance, shared terminology, and lifecycle processes, helping to integrate AI risks into enterprise programs and reduce confusion across teams is the correct answer. ISO/IEC 23894:2023 offers AI-specific risk management guidance, ISO/IEC 22989:2022 defines AI concepts and terminology, and ISO/IEC 5338:2023 outlines AI lifecycle processes and AI-specific concerns like concept drift and model decay. Together, they provide a common language, structured lifecycle, and risk discipline, making collaboration with auditors, vendors, and internal teams much easier.

Collectively establishing a single binding EU regulation that bans unacceptable-risk AI applications and imposes steep penalties for noncompliance across all Member States is incorrect. This describes the EU AI Act, not ISO standards. The EU AI Act is a law with risk tiers and penalties; ISO/IEC 23894, 22989, and 5338 are voluntary international standards that support, but do not replace or constitute, regulation. They do not impose legal penalties.

Functioning solely as consumer protection laws, focusing on deceptive AI advertising, and providing no guidance on technical lifecycle or risk processes is incorrect. These ISO standards are not consumer protection laws. They address terminology, lifecycle, and risk management for AI systems and are intended for technical, risk, and governance teams.

Providing primarily marketing-focused documents that standardize how AI vendors should describe product features and pricing tiers to enterprise buyers is incorrect. The standards are not marketing documents. Their focus is on conceptual clarity, lifecycle discipline, and risk management, enabling consistent engineering and governance practices rather than product packaging or pricing communication.

Related Content



6.3.2 Important AI Compliance Frameworks

resources\questions\q_important_ai_compliance_frameworks_03.question.xml

A mid-sized healthcare organization is rolling out a new AI-powered diagnostic tool across several medical departments. Leadership is concerned about maintaining oversight and regulatory compliance as the system uses sensitive patient data and learns from real-time inputs.

As part of establishing AI governance, your team is tasked with recommending technical guardrails that ensure secure and compliant AI operations.

Which of the following actions BEST exemplifies applying technical guardrails to support AI governance in this scenario?

- ☒ Implementing access controls that limit model deployment and sensitive data access to authorized personnel, along with automated logging and monitoring of every model inference and data transaction. ✓ Correct
- ☐ Conducting periodic staff training to raise awareness about ethical AI use and publishing a general AI policy document emphasizing patient data privacy.
- ☐ Using anonymized datasets during model development and reviewing the AI system after incidents, but allowing developers open access to production data for troubleshooting.
- ☐ Appointing a joint review committee from IT and compliance to sign off on new AI features before release, focusing on up-front risk approval.

Explanation

Implementing access controls that limit model deployment and sensitive data access to authorized personnel, along with automated logging and monitoring of every model inference and data transaction is the correct answer. This option directly implements technical guardrails: access controls based on least privilege, logging and audit trails, and real-time monitoring are explicit, technical measures embedded in the platform to enforce policy, ensure accountability, and maintain visibility over AI use. These measures support compliance, security, and the organization's governance framework.

Conducting periodic staff training to raise awareness about ethical AI use and publishing a general AI policy document is incorrect. They are not technical guardrails. Technical guardrails are programmatic, system-level controls that enforce governance automatically (such as access control, logging, and data masking). Policies and training support governance but do not directly enforce or monitor AI activity in the operational environment.

Appointing a joint review committee from IT and compliance to sign off on new AI features before release, focusing on up-front risk approval is incorrect. These not a technical guardrail. While important for governance, this approach primarily provides human review and may ensure policies are considered. However, it does not automatically and continuously enforce compliance or provide monitoring capabilities in the deployment and operation of AI systems.

Using anonymized datasets during model development and reviewing the AI system after incidents, but allowing developers open access to production data is incorrect. Allowing open developer access to production data contradicts technical guardrail principles such as "least privilege" and strict access controls. In fact, this approach increases the risk of data breaches, unauthorized changes, and reduces accountability. Further, it demonstrates a reactive rather than proactive control model.

Related Content

resources\questions\q_establish_ai_governance_08.question.xml

Why does the industry recommend treating model providers, data brokers, annotation vendors, and toolmakers as regulated suppliers?

- ☒ Because they should provide documentation, security attestations, bias-testing evidence, and commitments to update their offerings to support compliant AI use ✓ Correct
- ☐ Because outsourcing AI components to third parties automatically transfers all legal, security, and fairness responsibilities away from the organization
- ☐ Because regulated suppliers rarely change their products, so once approved, their tools can be used indefinitely without further review or monitoring
- ☐ Because relying on third-party tools allows the organization to skip its own AI governance processes, since the vendors have already implemented their own controls

Explanation

Because they should provide documentation, security attestations, bias-testing evidence, and commitments to update their offerings to support compliant AI use is the correct answer. Organizations should require model providers to provide documentation, security attestations, evidence of bias testing, and update commitments. This supports compliant, auditable AI use.

Because outsourcing AI components to third parties automatically transfers all legal, security, and fairness responsibilities away from the organization is incorrect. Responsibility cannot be fully transferred. Even when using third parties, the adopting organization remains accountable for how AI affects users, rights, and safety.

Because regulated suppliers rarely change their products, so once approved, their tools can be used indefinitely without further review or monitoring is incorrect. Third-party offerings evolve. It is important to focus on documentation, updates, and evidence implies ongoing evaluation, not one-time approval.

Because relying on third-party tools allows the organization to skip its own AI governance processes, since the vendors have already implemented their own controls, is incorrect. Third-party controls complement, but never replace, internal governance. The organization still needs its own policies, inventories, and monitoring.

Related Content

 6.3.1 Common Themes in AI Regulation

 6.3.5 External Compliance Impacts

resources\questions\q_common_themes_in_ai_regulation_07.question.xml

A retail company has developed a machine learning model to forecast product demand. The Data Engineer has delivered clean datasets, the Data Scientist has built an effective prototype, and the Machine Learning Engineer has optimized the model for production.

Now, the team needs to deploy the model in a way that automatically tracks which version is running in production, manages prompt and model version history, facilitates retraining with updated data, and allows seamless rollback to previous versions if needed.

Which role is primarily responsible for implementing these automated model deployment, versioning, retraining, and rollback capabilities within the AI platform?

- ☐ AI Security Architect
- ☒ MLOps Engineer ✓ Correct
- ☐ Platform Engineer
- ☐ Data Scientist

Explanation

MLOps Engineer is the correct answer. The MLOps Engineer is responsible for automating the deployment pipeline, managing model and prompt versioning, orchestrating scheduled retraining, and implementing monitoring and rollback processes to ensure the forecasting model operates reliably and can be quickly updated or reverted. The MLOps Engineer specializes in automating the end-to-end model lifecycle, including deployment, version control, retraining, and rollback mechanisms. This enables efficient updates, maintenance, and recovery for AI models in production.

Data Scientist is incorrect. While Data Scientists are vital for building and improving models, the responsibility for deployment automation, versioning, retraining orchestration, and rollback setup lies with the MLOps Engineer. Data Scientists focus on building effective prototypes and conducting analysis, not managing production pipelines.

Platform Engineer is incorrect. Platform Engineers maintain the foundational infrastructure required for AI development and deployment (such as compute and storage), but they are not responsible for automating model lifecycles or managing version histories, retraining, and rollbacks.

AI Security Architect is incorrect. The primary focus of the AI Security Architect is to safeguard models and data through security measures (such as access control and threat mitigation), not to automate deployment, versioning, retraining, or rollback processes.

Related Content



6.1.2 Important Roles in AI

resources\questions\q_important_roles_in_ai_05.question.xml